

final-ds-1

May 2, 2024

1 Data Wrangling, I

Perform the following operations using Python on any open source dataset (e.g., data.csv) 1. Import all the required Python Libraries. 2. Locate open source data from the web (e.g., <https://www.kaggle.com>). Provide a clear description of the data and its source (i.e., URL of the web site). 3. Load the Dataset into pandas dataframe. 4. Data Preprocessing: check for missing values in the data using pandas `isnull()`, `describe()` function to get some initial statistics. Provide variable descriptions. Types of variables etc. Check the dimensions of the data frame. 5. Data Formatting and Data Normalization: Summarize the types of variables by checking the data types (i.e., character, numeric, integer, factor, and logical) of the variables in the data set. If variables are not in the correct data type, apply proper type conversions. 6. Turn categorical variables into quantitative variables in Python.

In addition to the codes and outputs, explain every operation that you do in the above steps and explain everything that you do to import/read/scrape the data set.

```
[1]: import pandas as pd
```

```
[2]: df = pd.read_csv('StudentsPerformance.csv')
df
```

```
[2]:
```

| | gender | race/ethnicity | parental level of education | lunch | \ |
|-----|--------|----------------|-----------------------------|--------------|---|
| 0 | female | group B | bachelor's degree | standard | |
| 1 | female | group C | some college | standard | |
| 2 | female | group B | master's degree | standard | |
| 3 | male | group A | associate's degree | free/reduced | |
| 4 | male | group C | some college | standard | |
| .. | ... | ... | ... | ... | |
| 995 | female | group E | master's degree | standard | |
| 996 | male | group C | high school | free/reduced | |
| 997 | female | group C | high school | free/reduced | |
| 998 | female | group D | some college | standard | |
| 999 | female | group D | some college | free/reduced | |

| | test preparation course | math score | reading score | writing score |
|---|-------------------------|------------|---------------|---------------|
| 0 | none | 72 | 72 | 74 |
| 1 | completed | 69 | 90 | 88 |
| 2 | none | 90 | 95 | 93 |

| | | | | |
|-----|-----------|-----|-----|-----|
| 3 | none | 47 | 57 | 44 |
| 4 | none | 76 | 78 | 75 |
| .. | ... | ... | ... | ... |
| 995 | completed | 88 | 99 | 95 |
| 996 | none | 62 | 55 | 55 |
| 997 | completed | 59 | 71 | 65 |
| 998 | completed | 68 | 78 | 77 |
| 999 | none | 77 | 86 | 86 |

[1000 rows x 8 columns]

2 Date Preprocessing

```
[3]: df.isnull()
```

```
[3]:      gender  race/ethnicity  parental level of education  lunch  \
0      False      False      False      False  False
1      False      False      False      False  False
2      False      False      False      False  False
3      False      False      False      False  False
4      False      False      False      False  False
..      ...      ...      ...      ...      ...
995    False      False      False      False  False
996    False      False      False      False  False
997    False      False      False      False  False
998    False      False      False      False  False
999    False      False      False      False  False

      test preparation course  math score  reading score  writing score
0      False      False      False      False      False
1      False      False      False      False      False
2      False      False      False      False      False
3      False      False      False      False      False
4      False      False      False      False      False
..      ...      ...      ...      ...      ...
995    False      False      False      False      False
996    False      False      False      False      False
997    False      False      False      False      False
998    False      False      False      False      False
999    False      False      False      False      False
```

[1000 rows x 8 columns]

```
[4]: df.isnull().sum()
```

```
[4]: gender                0
     race/ethnicity        0
     parental level of education  0
     lunch                 0
     test preparation course  0
     math score            0
     reading score         0
     writing score          0
     dtype: int64
```

```
[5]: df.describe()
```

```
[5]:      math score  reading score  writing score
count  1000.00000    1000.000000    1000.000000
mean     66.08900      69.169000      68.054000
std      15.16308      14.600192      15.195657
min         0.00000      17.000000      10.000000
25%       57.00000      59.000000      57.750000
50%       66.00000      70.000000      69.000000
75%       77.00000      79.000000      79.000000
max      100.00000     100.000000     100.000000
```

```
[6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 8 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   gender                                1000 non-null   object
1   race/ethnicity                        1000 non-null   object
2   parental level of education           1000 non-null   object
3   lunch                                 1000 non-null   object
4   test preparation course               1000 non-null   object
5   math score                           1000 non-null   int64
6   reading score                        1000 non-null   int64
7   writing score                         1000 non-null   int64
dtypes: int64(3), object(5)
memory usage: 62.6+ KB
```

```
[7]: df.shape
```

```
[7]: (1000, 8)
```

```
[10]: df['writing score'] = df['writing score'].astype(float)
```

```
[11]: df
```

```
[11]:      gender race/ethnicity parental level of education      lunch \
0      female      group B      bachelor's degree      standard
1      female      group C      some college      standard
2      female      group B      master's degree      standard
3      male      group A      associate's degree free/reduced
4      male      group C      some college      standard
..      ...      ...      ...      ...
995     female      group E      master's degree      standard
996      male      group C      high school free/reduced
997     female      group C      high school free/reduced
998     female      group D      some college      standard
999     female      group D      some college free/reduced

      test preparation course math score reading score writing score
0      none      72      72      74.0
1      completed      69      90      88.0
2      none      90      95      93.0
3      none      47      57      44.0
4      none      76      78      75.0
..      ...      ...      ...      ...
995     completed      88      99      95.0
996      none      62      55      55.0
997     completed      59      71      65.0
998     completed      68      78      77.0
999      none      77      86      86.0

[1000 rows x 8 columns]
```

```
[14]: df['gender'].replace({'male' : 1 , 'female' : 0} , inplace=True)
```

```
[15]: df
```

```
[15]:      gender race/ethnicity parental level of education      lunch \
0          0      group B      bachelor's degree      standard
1          0      group C      some college      standard
2          0      group B      master's degree      standard
3          1      group A      associate's degree free/reduced
4          1      group C      some college      standard
..      ...      ...      ...      ...
995         0      group E      master's degree      standard
996         1      group C      high school free/reduced
997         0      group C      high school free/reduced
998         0      group D      some college      standard
999         0      group D      some college free/reduced

      test preparation course math score reading score writing score
0      none      72      72      74.0
```

| | | | | |
|-----|-----------|-----|-----|------|
| 1 | completed | 69 | 90 | 88.0 |
| 2 | none | 90 | 95 | 93.0 |
| 3 | none | 47 | 57 | 44.0 |
| 4 | none | 76 | 78 | 75.0 |
| .. | ... | ... | ... | ... |
| 995 | completed | 88 | 99 | 95.0 |
| 996 | none | 62 | 55 | 55.0 |
| 997 | completed | 59 | 71 | 65.0 |
| 998 | completed | 68 | 78 | 77.0 |
| 999 | none | 77 | 86 | 86.0 |

[1000 rows x 8 columns]

```
[16]: df.ndim
```

```
[16]: 2
```

3 To learn

1. Normalization and its techniques and way to import it in sklearn
2. One Hot encoding and Label Encoder

4 What is Data Preprocessing ?

Data preprocessing is the process of transforming raw data into a suitable format for analysis and machine learning. It involves cleaning the data (handling missing values, removing duplicates), transforming it (scaling, encoding categorical variables), and sometimes reducing its size (dimensionality reduction, feature selection).

This step is crucial to ensure models are trained on accurate, consistent, and standardized data, leading to better performance and more reliable results. Proper preprocessing simplifies data analysis and enhances interpretability.

5 What is One Hot Encoding?

One-hot encoding is a method used to convert categorical variables into a format suitable for machine learning models.

It creates new binary columns for each unique category in the original feature.

If a feature has three categories (“Red,” “Green,” “Blue”), one-hot encoding generates three new columns, with a value of “1” indicating the presence of that category and “0” otherwise.

6 What is Label Encoding ?

Label encoding is a method for converting categorical variables into numerical format by assigning a unique integer to each category.

If a feature has categories like “Small,” “Medium,” and “Large,” label encoding might map them to 0, 1, and 2, respectively.

7 Types of Scaling

Min-Max Scaling

Range between 0 and 1

Formula : $X' = \frac{X - \min(X)}{\max(X) - \min(X)}$

Standardization

Have a mean of 0 and a standard deviation of 1.

$X = \frac{X - \mu}{\sigma}$

[]:

final-ds-2

May 2, 2024

1 Data Wrangling II

Create an “Academic performance” dataset of students and perform the following operations using Python.

1. Scan all variables for missing values and inconsistencies. If there are missing values and/or inconsistencies, use any of the suitable techniques to deal with them.
2. Scan all numeric variables for outliers. If there are outliers, use any of the suitable techniques to deal with them.
3. Apply data transformations on at least one of the variables. The purpose of this transformation should be one of the following reasons: to change the scale for better understanding of the variable, to convert a non-linear relation into a linear one, or to decrease the skewness and convert the distribution into a normal distribution.

Reason and document your approach properly.

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import warnings

# Ignore all warnings
warnings.filterwarnings("ignore")
```

```
[2]: df = pd.read_csv('StudentsPerformance.csv')
df
```

```
[2]:
```

| | gender | race/ethnicity | parental level of education | lunch | \ |
|-----|--------|----------------|-----------------------------|--------------|---|
| 0 | female | group B | bachelor's degree | standard | |
| 1 | female | group C | some college | standard | |
| 2 | female | group B | master's degree | standard | |
| 3 | male | group A | associate's degree | free/reduced | |
| 4 | male | group C | some college | standard | |
| .. | ... | ... | ... | ... | |
| 995 | female | group E | master's degree | standard | |
| 996 | male | group C | high school | free/reduced | |
| 997 | female | group C | high school | free/reduced | |
| 998 | female | group D | some college | standard | |

| | | | | | |
|-----|-------------------------|------------|---------------|---------------|--------------|
| 999 | female | group D | | some college | free/reduced |
| | test preparation course | math score | reading score | writing score | |
| 0 | none | 72 | 72 | 74 | |
| 1 | completed | 69 | 90 | 88 | |
| 2 | none | 90 | 95 | 93 | |
| 3 | none | 47 | 57 | 44 | |
| 4 | none | 76 | 78 | 75 | |
| .. | ... | ... | ... | ... | |
| 995 | completed | 88 | 99 | 95 | |
| 996 | none | 62 | 55 | 55 | |
| 997 | completed | 59 | 71 | 65 | |
| 998 | completed | 68 | 78 | 77 | |
| 999 | none | 77 | 86 | 86 | |

[1000 rows x 8 columns]


1. Scan all variables for missing values and inconsistencies. If there are missing values and/or inconsistencies, use any of the suitable techniques to deal with them.

for this use fillna() method either by filling it with mean , mode , median

df['col'] = df['col'].fillna(df['col'].mean() , inplace = True)

```
[3]: df.boxplot()
```

```
[3]: <Axes: >
```

```
[4]: Q1 = df['math score'].quantile(0.25)
Q3 = df['math score'].quantile(0.75)
IQR = Q3 - Q1
lower_bound = Q1 - (1.5 * IQR)
upper_bound = Q3 + (1.5 * IQR)

outliers = df[(df['math score'] < lower_bound) | (df['math score'] >
↪upper_bound)]
outliers
```

```
[4]:      gender race/ethnicity parental level of education      lunch \
17   female      group B      some high school  free/reduced
59   female      group C      some high school  free/reduced
145  female      group C      some college    free/reduced
338  female      group B      some high school  free/reduced
466  female      group D  associate's degree  free/reduced
787  female      group B      some college    standard
842  female      group B      high school    free/reduced
980  female      group B      high school    free/reduced

      test preparation course  math score  reading score  writing score
17              none          18          32          28
```

| | | | | |
|-----|-----------|----|----|----|
| 59 | none | 0 | 17 | 10 |
| 145 | none | 22 | 39 | 33 |
| 338 | none | 24 | 38 | 27 |
| 466 | none | 26 | 31 | 38 |
| 787 | none | 19 | 38 | 32 |
| 842 | completed | 23 | 44 | 36 |
| 980 | none | 8 | 24 | 23 |

```
[5]: df_cleaned = df[(df['math score'] >= lower_bound) & (df['math score'] <=
    ↪upper_bound)]
```

```
[6]: df_cleaned
```


```
[6]:   gender race/ethnicity parental level of education      lunch \
0   female      group B      bachelor's degree      standard
1   female      group C      some college      standard
2   female      group B      master's degree      standard
3    male      group A      associate's degree  free/reduced
4    male      group C      some college      standard
..   ...      ...      ...      ...
995  female      group E      master's degree      standard
996   male      group C      high school  free/reduced
997  female      group C      high school  free/reduced
998  female      group D      some college      standard
999  female      group D      some college  free/reduced

      test preparation course  math score  reading score  writing score
0              none          72          72          74
1      completed          69          90          88
2              none          90          95          93
3              none          47          57          44
4              none          76          78          75
..   ...      ...      ...      ...
995      completed          88          99          95
996              none          62          55          55
997      completed          59          71          65
998      completed          68          78          77
999              none          77          86          86
```

[992 rows x 8 columns]

```
[7]: df_cleaned.boxplot()
```

```
[7]: <Axes: >
```



2 Transformation

```
[8]: df_cleaned['gender'] = df_cleaned['gender'].replace({'male' : 1 , 'female' : 0})
```

```
[9]: df_cleaned
```

```
[9]:
```

| | gender | race/ethnicity | parental level of education | lunch | \ |
|-----|--------|----------------|-----------------------------|--------------|---|
| 0 | 0 | group B | bachelor's degree | standard | |
| 1 | 0 | group C | some college | standard | |
| 2 | 0 | group B | master's degree | standard | |
| 3 | 1 | group A | associate's degree | free/reduced | |
| 4 | 1 | group C | some college | standard | |
| .. | ... | ... | ... | ... | |
| 995 | 0 | group E | master's degree | standard | |
| 996 | 1 | group C | high school | free/reduced | |
| 997 | 0 | group C | high school | free/reduced | |
| 998 | 0 | group D | some college | standard | |
| 999 | 0 | group D | some college | free/reduced | |


| | test preparation course | math score | reading score | writing score |
|---|-------------------------|------------|---------------|---------------|
| 0 | none | 72 | 72 | 74 |

| | | | | |
|-----|-----------|-----|-----|-----|
| 1 | completed | 69 | 90 | 88 |
| 2 | none | 90 | 95 | 93 |
| 3 | none | 47 | 57 | 44 |
| 4 | none | 76 | 78 | 75 |
| .. | ... | ... | ... | ... |
| 995 | completed | 88 | 99 | 95 |
| 996 | none | 62 | 55 | 55 |
| 997 | completed | 59 | 71 | 65 |
| 998 | completed | 68 | 78 | 77 |
| 999 | none | 77 | 86 | 86 |

[992 rows x 8 columns]

[]:

```
[10]: fig, axes= plt.subplots(1,2)
sns.boxplot(df['math score'],ax=axes[0])
axes[0].title.set_text('Before')
sns.boxplot(df_cleaned['math score'],ax=axes[1])
axes[1].title.set_text('After')
plt.show()
```



```
[11]: from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()
df_cleaned['reading score'] = scaler.fit_transform(df_cleaned[['reading_
↳score']])

df_cleaned
```

```
[11]:      gender race/ethnicity parental level of education      lunch \
0         0      group B      bachelor's degree      standard
1         0      group C      some college      standard
2         0      group B      master's degree      standard
3         1      group A      associate's degree free/reduced
4         1      group C      some college      standard
..      ...      ...      ...      ...
995        0      group E      master's degree      standard
996        1      group C      high school free/reduced
997        0      group C      high school free/reduced
998        0      group D      some college      standard
999        0      group D      some college free/reduced

      test preparation course  math score  reading score  writing score
0              none          72      0.636364          74
1      completed          69      0.870130          88
2              none          90      0.935065          93
3              none          47      0.441558          44
4              none          76      0.714286          75
..      ...      ...      ...      ...
995      completed          88      0.987013          95
996              none          62      0.415584          55
997      completed          59      0.623377          65
998      completed          68      0.714286          77
999              none          77      0.818182          86
```

[992 rows x 8 columns]

```
[ ]:
```

final-ds-3

May 2, 2024

1 Descriptive Statistics - Measures of Central Tendency and variability

Perform the following operations on any open source dataset (e.g., data.csv) 1. Provide summary statistics (mean, median, minimum, maximum, standard deviation) for a dataset (age, income etc.) with numeric variables grouped by one of the qualitative (categorical) variable. For example, if your categorical variable is age groups and quantitative variable is income, then provide summary statistics of income grouped by the age groups. Create a list that contains a numeric value for each response to the categorical variable. 2. Write a Python program to display some basic statistical details like percentile, mean, standard deviation etc. of the species of 'Iris-setosa', 'Iris-versicolor' and 'Iris-versicolor' of iris.csv dataset. Provide the codes with outputs and explain everything that you do in this step

```
[1]: import pandas as pd
```

```
[2]: df = pd.read_csv('academic_performance.csv')
df
```

```
[2]:
```

| | Student_ID | Gender | Age | Math_Score | Science_Score | English_Score |
|-----|------------|--------|-----|------------|---------------|---------------|
| 0 | S0001 | Male | 23 | 19 | 71 | 90 |
| 1 | S0002 | Female | 18 | 22 | 25 | 15 |
| 2 | S0003 | Female | 23 | 11 | 100 | 68 |
| 3 | S0004 | Male | 17 | 11 | 77 | 74 |
| 4 | S0005 | Female | 22 | 33 | 50 | 85 |
| .. | ... | ... | ... | ... | ... | ... |
| 495 | S0496 | Male | 17 | 65 | 96 | 51 |
| 496 | S0497 | Female | 16 | 99 | 61 | 34 |
| 497 | S0498 | Male | 16 | 71 | 1 | 58 |
| 498 | S0499 | Female | 21 | 1 | 47 | 80 |
| 499 | S0500 | Male | 22 | 1 | 43 | 28 |

[500 rows x 6 columns]

```
[3]: df.describe()
```

```
[3]:
```

| | Age | Math_Score | Science_Score | English_Score |
|-------|------------|------------|---------------|---------------|
| count | 500.000000 | 500.000000 | 500.000000 | 500.000000 |
| mean | 20.438000 | 49.494000 | 51.040000 | 50.188000 |

| | | | | |
|-----|-----------|------------|------------|------------|
| std | 2.833056 | 30.141782 | 29.182969 | 28.702968 |
| min | 16.000000 | 0.000000 | 0.000000 | 0.000000 |
| 25% | 18.000000 | 22.750000 | 25.000000 | 26.750000 |
| 50% | 20.000000 | 49.000000 | 51.000000 | 51.000000 |
| 75% | 23.000000 | 77.000000 | 76.250000 | 74.000000 |
| max | 25.000000 | 100.000000 | 100.000000 | 100.000000 |

```
[4]: df.groupby('Gender')[['Math_Score' , 'Science_Score' , 'English_Score']].sum()
```

```
[4]:      Math_Score  Science_Score  English_Score
Gender
Female      12287           13260           13070
Male       12460           12260           12024
```

```
[5]: df.groupby('Gender')[['Math_Score' , 'Science_Score' , 'English_Score']].
      describe().T
```

```
[5]: Gender
Math_Score  count  Female  Male
          mean    46.896947  52.352941
          std     29.865485  30.248224
          min      0.000000   1.000000
          25%     20.000000  24.250000
          50%     44.000000  56.000000
          75%     73.750000  81.000000
          max    100.000000 100.000000
Science_Score count  Female  Male
          mean    50.610687  51.512605
          std     28.881200  29.565308
          min      1.000000   0.000000
          25%     25.250000  25.000000
          50%     51.000000  50.500000
          75%     75.000000  78.000000
          max    100.000000 100.000000
English_Score count  Female  Male
          mean    49.885496  50.521008
          std     28.364119  29.127640
          min      0.000000   0.000000
          25%     27.000000  26.000000
          50%     50.000000  53.000000
          75%     73.500000  74.750000
          max    100.000000 100.000000
```

```
[6]: flower = pd.read_csv('Iris (1).csv')
flower
```

```
[6]:      Id  SepalLengthCm  SepalWidthCm  PetalLengthCm  PetalWidthCm  \
0      1          5.1          3.5          1.4          0.2
1      2          4.9          3.0          1.4          0.2
2      3          4.7          3.2          1.3          0.2
3      4          4.6          3.1          1.5          0.2
4      5          5.0          3.6          1.4          0.2
..    ...
145   146          6.7          3.0          5.2          2.3
146   147          6.3          2.5          5.0          1.9
147   148          6.5          3.0          5.2          2.0
148   149          6.2          3.4          5.4          2.3
149   150          5.9          3.0          5.1          1.8
```

```
      Species
0      Iris-setosa
1      Iris-setosa
2      Iris-setosa
3      Iris-setosa
4      Iris-setosa
..    ...
145   Iris-virginica
146   Iris-virginica
147   Iris-virginica
148   Iris-virginica
149   Iris-virginica
```

[150 rows x 6 columns]

```
[7]: flower.describe()
```

```
[7]:      Id  SepalLengthCm  SepalWidthCm  PetalLengthCm  PetalWidthCm
count  150.000000      150.000000      150.000000      150.000000      150.000000
mean    75.500000         5.843333         3.054000         3.758667         1.198667
std     43.445368         0.828066         0.433594         1.764420         0.763161
min       1.000000         4.300000         2.000000         1.000000         0.100000
25%     38.250000         5.100000         2.800000         1.600000         0.300000
50%     75.500000         5.800000         3.000000         4.350000         1.300000
75%    112.750000         6.400000         3.300000         5.100000         1.800000
max    150.000000         7.900000         4.400000         6.900000         2.500000
```

```
[8]: f = flower.groupby(by='Species')
```

```
[9]: f.first()
```

```
[9]:      Id  SepalLengthCm  SepalWidthCm  PetalLengthCm  PetalWidthCm
Species
Iris-setosa      1          5.1          3.5          1.4          0.2
```


| | | | | | |
|-----------------|-----|-----|-----|-----|-----|
| Iris-versicolor | 51 | 7.0 | 3.2 | 4.7 | 1.4 |
| Iris-virginica | 101 | 6.3 | 3.3 | 6.0 | 2.5 |

```
[10]: flower.groupby('Species').describe().T
```

```
[10]: Species
Id      count  50.000000  50.000000  50.000000
      mean    25.500000  75.500000  125.500000
      std     14.577380  14.577380  14.577380
      min      1.000000  51.000000  101.000000
      25%     13.250000  63.250000  113.250000
      50%     25.500000  75.500000  125.500000
      75%     37.750000  87.750000  137.750000
      max     50.000000  100.000000  150.000000
SepalLengthCm count  50.000000  50.000000  50.000000
      mean      5.006000   5.936000   6.588000
      std      0.352490   0.516171   0.635880
      min      4.300000   4.900000   4.900000
      25%      4.800000   5.600000   6.225000
      50%      5.000000   5.900000   6.500000
      75%      5.200000   6.300000   6.900000
      max      5.800000   7.000000   7.900000
SepalWidthCm  count  50.000000  50.000000  50.000000
      mean      3.418000   2.770000   2.974000
      std      0.381024   0.313798   0.322497
      min      2.300000   2.000000   2.200000
      25%      3.125000   2.525000   2.800000
      50%      3.400000   2.800000   3.000000
      75%      3.675000   3.000000   3.175000
      max      4.400000   3.400000   3.800000
PetalLengthCm count  50.000000  50.000000  50.000000
      mean      1.464000   4.260000   5.552000
      std      0.173511   0.469911   0.551895
      min      1.000000   3.000000   4.500000
      25%      1.400000   4.000000   5.100000
      50%      1.500000   4.350000   5.550000
      75%      1.575000   4.600000   5.875000
      max      1.900000   5.100000   6.900000
PetalWidthCm  count  50.000000  50.000000  50.000000
      mean      0.244000   1.326000   2.026000
      std      0.107210   0.197753   0.274650
      min      0.100000   1.000000   1.400000
      25%      0.200000   1.200000   1.800000
      50%      0.200000   1.300000   2.000000
      75%      0.300000   1.500000   2.300000
      max      0.600000   1.800000   2.500000
```

2 What is groupby in pandas ?

Groupby in pandas is a method that groups a DataFrame's rows based on one or more columns, allowing you to perform aggregate operations on each group.

3 What is Data Science ?

Data Science is an interdisciplinary field that combines statistical analysis, machine learning, data engineering, and domain expertise to extract insights and knowledge from data. It involves collecting, cleaning, analyzing, and interpreting data to inform decision-making and create data-driven solutions. Data Science is used across industries to solve complex problems and guide strategic decisions.

4 Pandas

Pandas is an open-source data analysis and manipulation library for Python, designed to work with structured data.

It provides powerful data structures like DataFrames and Series, allowing you to perform complex operations such as filtering, grouping, merging, and aggregating data with ease.

final-ds-4

May 2, 2024

1 Data Analytics I

Create a Linear Regression Model using Python/R to predict home prices using Boston Housing Dataset (<https://www.kaggle.com/c/boston-housing>). The Boston Housing dataset contains information about various houses in Boston through different parameters. There are 506 samples and 14 feature variables in this dataset.

The objective is to predict the value of prices of the house using the given features.

```
[1]: import pandas as pd
      from sklearn.linear_model import LinearRegression
      from sklearn.model_selection import train_test_split
      from sklearn.metrics import r2_score , mean_squared_error
      import numpy as np
```

```
[2]: df = pd.read_csv('boston-housing.csv')
      df
```

```
[2]:
```

| | ID | crim | zn | indus | chas | nox | rm | age | dis | rad | tax | \ |
|-----|-----|---------|------|-------|------|-------|-------|------|--------|-----|-----|-----|
| 0 | 1 | 0.00632 | 18.0 | 2.31 | 0 | 0.538 | 6.575 | 65.2 | 4.0900 | 1 | 296 | |
| 1 | 2 | 0.02731 | 0.0 | 7.07 | 0 | 0.469 | 6.421 | 78.9 | 4.9671 | 2 | 242 | |
| 2 | 4 | 0.03237 | 0.0 | 2.18 | 0 | 0.458 | 6.998 | 45.8 | 6.0622 | 3 | 222 | |
| 3 | 5 | 0.06905 | 0.0 | 2.18 | 0 | 0.458 | 7.147 | 54.2 | 6.0622 | 3 | 222 | |
| 4 | 7 | 0.08829 | 12.5 | 7.87 | 0 | 0.524 | 6.012 | 66.6 | 5.5605 | 5 | 311 | |
| .. | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 328 | 500 | 0.17783 | 0.0 | 9.69 | 0 | 0.585 | 5.569 | 73.5 | 2.3999 | 6 | 391 | |
| 329 | 502 | 0.06263 | 0.0 | 11.93 | 0 | 0.573 | 6.593 | 69.1 | 2.4786 | 1 | 273 | |
| 330 | 503 | 0.04527 | 0.0 | 11.93 | 0 | 0.573 | 6.120 | 76.7 | 2.2875 | 1 | 273 | |
| 331 | 504 | 0.06076 | 0.0 | 11.93 | 0 | 0.573 | 6.976 | 91.0 | 2.1675 | 1 | 273 | |
| 332 | 506 | 0.04741 | 0.0 | 11.93 | 0 | 0.573 | 6.030 | 80.8 | 2.5050 | 1 | 273 | |

| | ptratio | black | lstat | medv |
|----|---------|--------|-------|------|
| 0 | 15.3 | 396.90 | 4.98 | 24.0 |
| 1 | 17.8 | 396.90 | 9.14 | 21.6 |
| 2 | 18.7 | 394.63 | 2.94 | 33.4 |
| 3 | 18.7 | 396.90 | 5.33 | 36.2 |
| 4 | 15.2 | 395.60 | 12.43 | 22.9 |
| .. | ... | ... | ... | ... |

```

328    19.2  395.77  15.10  17.5
329    21.0  391.99   9.67  22.4
330    21.0  396.90   9.08  20.6
331    21.0  396.90   5.64  23.9
332    21.0  396.90   7.88  11.9

```

```
[333 rows x 15 columns]
```

```
[3]: df.describe()
```

```

[3]:
count    ID      crim      zn      indus      chas      nox  \
count  333.000000  333.000000  333.000000  333.000000  333.000000  333.000000
mean   250.951952   3.360341  10.689189   11.293483   0.060060   0.557144
std    147.859438   7.352272  22.674762   6.998123   0.237956   0.114955
min      1.000000   0.006320   0.000000   0.740000   0.000000   0.385000
25%    123.000000   0.078960   0.000000   5.130000   0.000000   0.453000
50%    244.000000   0.261690   0.000000   9.900000   0.000000   0.538000
75%    377.000000   3.678220  12.500000  18.100000   0.000000   0.631000
max    506.000000  73.534100  100.000000  27.740000   1.000000   0.871000

count    rm      age      dis      rad      tax      ptratio  \
count  333.000000  333.000000  333.000000  333.000000  333.000000  333.000000
mean     6.265619  68.226426   3.709934   9.633634  409.279279  18.448048
std     0.703952  28.133344   1.981123   8.742174  170.841988   2.151821
min     3.561000   6.000000   1.129600   1.000000  188.000000  12.600000
25%     5.884000  45.400000   2.122400   4.000000  279.000000  17.400000
50%     6.202000  76.700000   3.092300   5.000000  330.000000  19.000000
75%     6.595000  93.800000   5.116700  24.000000  666.000000  20.200000
max     8.725000 100.000000  10.710300  24.000000  711.000000  21.200000

count    black    lstat    medv
count  333.000000  333.000000  333.000000
mean   359.466096  12.515435  22.768769
std    86.584567   7.067781   9.173468
min     3.500000   1.730000   5.000000
25%    376.730000   7.180000  17.400000
50%    392.050000  10.970000  21.600000
75%    396.240000  16.420000  25.000000
max    396.900000  37.970000  50.000000

```

```
[4]: df.columns
```

```

[4]: Index(['ID', 'crim', 'zn', 'indus', 'chas', 'nox', 'rm', 'age', 'dis', 'rad',
          'tax', 'ptratio', 'black', 'lstat', 'medv'],
          dtype='object')

```

```
[5]: x = df[['ID', 'crim', 'zn', 'indus', 'chas', 'nox', 'rm', 'age', 'dis',  
↳ 'rad', 'tax', 'ptratio', 'black', 'lstat']]  
y = df['medv']
```

```
[6]: x_train , x_test , y_train , y_test = train_test_split(x,y,test_size=0.25 ,  
↳ random_state=42)
```

```
[7]: lr = LinearRegression()
```

```
[8]: lr.fit(x_train , y_train)
```

```
[8]: LinearRegression()
```

```
[9]: y_pred = lr.predict(x_test)  
y_pred
```

```
[9]: array([25.42935235, 22.96355416, 22.98361313, 32.78417799, 25.00393879,  
14.02939811, 17.33416663, 30.26235415, 15.66094194, 25.17645124,  
26.57561017, 19.87047671, 20.0909698 , 34.65458829, 21.54826488,  
34.56267042, 22.2485164 , 16.882139 , 25.36436429, 16.87727359,  
36.17705665, 31.39060697, 22.8003689 , 28.13000873, 17.20049765,  
42.37151734, 7.87497041, -0.36689721, 31.09480595, 8.4983594 ,  
19.11933223, 20.39461378, 27.46125616, 15.1123262 , 19.49023076,  
12.24379213, 27.56523215, 4.49445516, 17.2050638 , 22.53286675,  
24.31038054, 22.18772176, 25.17245639, 39.17706326, 36.30722879,  
21.36801908, 11.07627164, 21.45992975, 13.86557089, 20.5226888 ,  
13.18677931, 27.99316461, 21.49698121, 13.63192997, 33.57300717,  
1.31418106, 20.80538627, 27.16216548, 25.27755635, 28.15722941,  
17.87759693, 23.94807788, 18.11868011, 28.38039855, 23.49353731,  
16.58474878, 26.43944232, 20.85487585, 23.73111221, 16.67730384,  
14.83654746, 26.14692568, 21.12274293, 12.8361667 , 24.2684098 ,  
33.5090989 , 22.02205101, 23.54787016, 22.92262304, 16.50548743,  
21.16265595, 28.70685258, 40.24134644, 15.29564422])
```

```
[10]: lr.score(x_train , y_train) #Training Accuracy
```

```
[10]: 0.7322764285677805
```

```
[11]: lr.score(x_test,y_test) #Testing Accuracy
```

```
[11]: 0.7257587357992887
```

```
[12]: r2_Score = r2_score(y_test,y_pred) #closer to 1 indicates better model  
r2_Score
```

```
[12]: 0.7257587357992887
```


```
[13]: np.sqrt(mean_squared_error(y_test,y_pred)) # smaller rmse indicates better
      ↪ model
```

```
[13]: 4.823701046478235
```

```
[14]: import matplotlib.pyplot as plt
```

```
[15]: coef = np.polyfit(y_test,y_pred,1)
      poly = np.poly1d(coef)
      y_fit = poly(y_test)

      plt.scatter(x=y_test , y=y_pred)
      plt.plot(y_test,y_fit,color = 'red')
      plt.xlabel('True Values')
      plt.ylabel('Predictions')
      plt.show()
```



2 What is Linear Regression ?

Linear regression is a statistical method for modeling the relationship between a dependent variable (target) and one or more independent variables (predictors) using a linear equation. The goal is to

find the best-fitting line or hyperplane that minimizes the sum of the squared differences between observed and predicted values.

3 What is r2 Score ?

The R^2 score, or coefficient of determination, measures how well a regression model explains the variance in the dependent variable. It ranges from 0 to 1, where a score of 1 indicates perfect prediction, and a score of 0 means the model does not explain any variance beyond the mean.

[]:

final-ds-5

May 2, 2024

1 5) Data Analytics II

1. Implement logistic regression using Python/R to perform classification on Social_Network_Ads.csv dataset.
2. Compute Confusion matrix to find TP, FP, TN, FN, Accuracy, Error rate, Precision, Recall on the given dataset.

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score , confusion_matrix , precision_score,
    recall_score , classification_report
```

```
[2]: df = pd.read_csv('Social_Network_Ads.csv')
df
```

```
[2]:
```

| | User ID | Gender | Age | EstimatedSalary | Purchased |
|-----|----------|--------|-----|-----------------|-----------|
| 0 | 15624510 | Male | 19 | 19000 | 0 |
| 1 | 15810944 | Male | 35 | 20000 | 0 |
| 2 | 15668575 | Female | 26 | 43000 | 0 |
| 3 | 15603246 | Female | 27 | 57000 | 0 |
| 4 | 15804002 | Male | 19 | 76000 | 0 |
| .. | ... | ... | ... | ... | ... |
| 395 | 15691863 | Female | 46 | 41000 | 1 |
| 396 | 15706071 | Male | 51 | 23000 | 1 |
| 397 | 15654296 | Female | 50 | 20000 | 1 |
| 398 | 15755018 | Male | 36 | 33000 | 0 |
| 399 | 15594041 | Female | 49 | 36000 | 1 |

[400 rows x 5 columns]

```
[3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
```


| # | Column | Non-Null Count | Dtype |
|---|-----------------|----------------|--------|
| 0 | User ID | 400 non-null | int64 |
| 1 | Gender | 400 non-null | object |
| 2 | Age | 400 non-null | int64 |
| 3 | EstimatedSalary | 400 non-null | int64 |
| 4 | Purchased | 400 non-null | int64 |

dtypes: int64(4), object(1)
memory usage: 15.8+ KB

```
[4]: df.columns
```

```
[4]: Index(['User ID', 'Gender', 'Age', 'EstimatedSalary', 'Purchased'],
dtype='object')
```

```
[5]: # Male : 1 Female : 0
df['Gender'].replace({'Male' : 1 , 'Female' : 0} , inplace=True)
```

```
[6]: df
```

```
[6]:
```

| | User ID | Gender | Age | EstimatedSalary | Purchased |
|-----|----------|--------|-----|-----------------|-----------|
| 0 | 15624510 | 1 | 19 | 19000 | 0 |
| 1 | 15810944 | 1 | 35 | 20000 | 0 |
| 2 | 15668575 | 0 | 26 | 43000 | 0 |
| 3 | 15603246 | 0 | 27 | 57000 | 0 |
| 4 | 15804002 | 1 | 19 | 76000 | 0 |
| .. | ... | ... | ... | ... | ... |
| 395 | 15691863 | 0 | 46 | 41000 | 1 |
| 396 | 15706071 | 1 | 51 | 23000 | 1 |
| 397 | 15654296 | 0 | 50 | 20000 | 1 |
| 398 | 15755018 | 1 | 36 | 33000 | 0 |
| 399 | 15594041 | 0 | 49 | 36000 | 1 |

[400 rows x 5 columns]

```
[7]: x = df[['Gender', 'Age', 'EstimatedSalary']]
x
```

```
[7]:
```

| | Gender | Age | EstimatedSalary |
|-----|--------|-----|-----------------|
| 0 | 1 | 19 | 19000 |
| 1 | 1 | 35 | 20000 |
| 2 | 0 | 26 | 43000 |
| 3 | 0 | 27 | 57000 |
| 4 | 1 | 19 | 76000 |
| .. | ... | ... | ... |
| 395 | 0 | 46 | 41000 |
| 396 | 1 | 51 | 23000 |

```
397      0   50      20000
398      1   36      33000
399      0   49      36000
```

```
[400 rows x 3 columns]
```

```
[8]: y = df[['Purchased']]
      y
```

```
[8]:      Purchased
0           0
1           0
2           0
3           0
4           0
..         ...
395          1
396          1
397          1
398          0
399          1
```

```
[400 rows x 1 columns]
```

```
[9]: X_train, X_test, y_train, y_test = train_test_split(x,y,test_size=0.
      ↪25,random_state=42)
```

```
[10]: lr = LogisticRegression()
```

```
[11]: lr.fit(X_train,y_train)
```

```
C:\Users\vedan\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.9_qbz5n2
kfra8p0\LocalCache\local-packages\Python39\site-
packages\sklearn\utils\validation.py:1300: DataConversionWarning: A column-
vector y was passed when a 1d array was expected. Please change the shape of y
to (n_samples, ), for example using ravel().
```

```
    y = column_or_1d(y, warn=True)
```

```
C:\Users\vedan\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.9_qbz5n2
kfra8p0\LocalCache\local-packages\Python39\site-
packages\sklearn\linear_model\_logistic.py:469: ConvergenceWarning: lbfgs failed
to converge (status=1):
```

```
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-

```

regression
    n_iter_i = _check_optimize_result(
[11]: LogisticRegression()

[12]: y_pred = lr.predict(X_test)
      y_pred

[12]: array([0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0,
            0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0,
            0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0,
            1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1,
            0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0], dtype=int64)

[13]: accuracy = accuracy_score(y_test , y_pred)
      accuracy

[13]: 0.89

[27]: error_rate = 1 - accuracy
      round(error_rate,2)

[27]: 0.11

```

2 Confusion Matrix

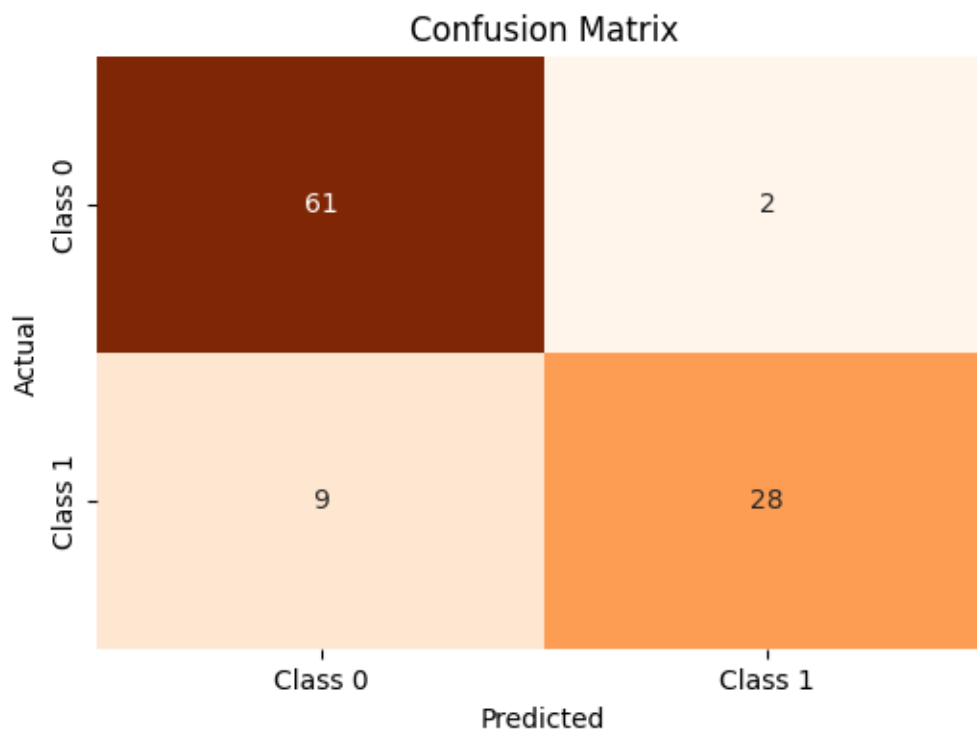
```

[14]: cm = confusion_matrix(y_test,y_pred)
      cm

[14]: array([[61,  2],
            [ 9, 28]], dtype=int64)

[15]: # Create a heatmap to visualize the confusion matrix
      plt.figure(figsize=(6, 4))
      sns.heatmap(cm, annot=True, fmt='d', cmap='Oranges', cbar=False,
                  xticklabels=['Class 0', 'Class 1'],
                  yticklabels=['Class 0', 'Class 1'])
      plt.xlabel('Predicted')
      plt.ylabel('Actual')
      plt.title('Confusion Matrix')
      plt.show()

```



```
[16]: # Sequence TN , FP , FN ,
```

```
[17]: #      Predicted
#      0    1
# Actual 0  TN  FP
#      1  FN  TP
```

```
[18]: precision = precision_score(y_test,y_pred)
precision
```

```
[18]: 0.9333333333333333
```

```
[19]: recall = recall_score(y_test,y_pred)
recall
```

```
[19]: 0.7567567567567568
```

```
[20]: report = classification_report(y_test,y_pred)
print(report)
```

```

          precision    recall  f1-score   support

0               0.87         0.97         0.92         63

```

| | | | | |
|--------------|------|------|------|-----|
| 1 | 0.93 | 0.76 | 0.84 | 37 |
| accuracy | | | 0.89 | 100 |
| macro avg | 0.90 | 0.86 | 0.88 | 100 |
| weighted avg | 0.89 | 0.89 | 0.89 | 100 |

3 See Formulae for Precision Recall

4 Precision

Precision measures the accuracy of positive predictions, calculated as the ratio of true positives to the total predicted positives. A high precision indicates that most of the positive predictions are correct. $\text{Precision} = \text{TP} / (\text{TP} + \text{FP})$

5 Recall

measures the proportion of actual positives that are correctly identified. It is calculated as the ratio of true positives to the total actual positives. $\text{Recall} = \text{TP} / (\text{TP} + \text{FN})$

6 True Positive Rate (TPR)

indicates how well a model identifies positive instances.

Formula same as Recall

7 True Negative Rate (TNR)

indicates how well a model avoids false positives.

$$\text{FPR} = \text{FP} / (\text{FP} + \text{TN})$$

[]:

final-ds-6

May 2, 2024

1 6) Data Analytics III

1. Implement Simple Naïve Bayes classification algorithm using Python/R on iris.csv dataset.
2. Compute Confusion matrix to find TP, FP, TN, FN, Accuracy, Error rate, Precision, Recall on the given dataset.

```
[1]: import pandas as pd
import numpy as np
from sklearn.naive_bayes import GaussianNB
```

```
[2]: df = pd.read_csv('Iris (1).csv')
df
```

```
[2]:
```

| | Id | SepalLengthCm | SepalWidthCm | PetalLengthCm | PetalWidthCm | \ |
|-----|-----|---------------|--------------|---------------|--------------|---|
| 0 | 1 | 5.1 | 3.5 | 1.4 | 0.2 | |
| 1 | 2 | 4.9 | 3.0 | 1.4 | 0.2 | |
| 2 | 3 | 4.7 | 3.2 | 1.3 | 0.2 | |
| 3 | 4 | 4.6 | 3.1 | 1.5 | 0.2 | |
| 4 | 5 | 5.0 | 3.6 | 1.4 | 0.2 | |
| .. | ... | ... | ... | ... | ... | |
| 145 | 146 | 6.7 | 3.0 | 5.2 | 2.3 | |
| 146 | 147 | 6.3 | 2.5 | 5.0 | 1.9 | |
| 147 | 148 | 6.5 | 3.0 | 5.2 | 2.0 | |
| 148 | 149 | 6.2 | 3.4 | 5.4 | 2.3 | |
| 149 | 150 | 5.9 | 3.0 | 5.1 | 1.8 | |

| | Species |
|-----|----------------|
| 0 | Iris-setosa |
| 1 | Iris-setosa |
| 2 | Iris-setosa |
| 3 | Iris-setosa |
| 4 | Iris-setosa |
| .. | ... |
| 145 | Iris-virginica |
| 146 | Iris-virginica |
| 147 | Iris-virginica |
| 148 | Iris-virginica |
| 149 | Iris-virginica |

[150 rows x 6 columns]

```
[3]: x = df.drop(columns=['Id' , 'Species'])
x
```

```
[3]:
```

| | SepalLengthCm | SepalWidthCm | PetalLengthCm | PetalWidthCm |
|-----|---------------|--------------|---------------|--------------|
| 0 | 5.1 | 3.5 | 1.4 | 0.2 |
| 1 | 4.9 | 3.0 | 1.4 | 0.2 |
| 2 | 4.7 | 3.2 | 1.3 | 0.2 |
| 3 | 4.6 | 3.1 | 1.5 | 0.2 |
| 4 | 5.0 | 3.6 | 1.4 | 0.2 |
| .. | ... | ... | ... | ... |
| 145 | 6.7 | 3.0 | 5.2 | 2.3 |
| 146 | 6.3 | 2.5 | 5.0 | 1.9 |
| 147 | 6.5 | 3.0 | 5.2 | 2.0 |
| 148 | 6.2 | 3.4 | 5.4 | 2.3 |
| 149 | 5.9 | 3.0 | 5.1 | 1.8 |

[150 rows x 4 columns]

```
[4]: y = df[['Species']]
y
```

```
[4]:
```

| | Species |
|-----|----------------|
| 0 | Iris-setosa |
| 1 | Iris-setosa |
| 2 | Iris-setosa |
| 3 | Iris-setosa |
| 4 | Iris-setosa |
| .. | ... |
| 145 | Iris-virginica |
| 146 | Iris-virginica |
| 147 | Iris-virginica |
| 148 | Iris-virginica |
| 149 | Iris-virginica |

[150 rows x 1 columns]

```
[5]: # from sklearn.preprocessing import MinMaxScaler

# scaler = MinMaxScaler()
```

```
[7]: # X_scaled = scaler.fit_transform(x)
# X_scaled
```

```
[8]: from sklearn.model_selection import train_test_split
```

```
[9]: X_train, X_test, y_train, y_test = train_test_split(x , y , test_size=0.25 ,  
↳ random_state=42)
```

```
[10]: gnb = GaussianNB()
```

```
[11]: gnb.fit(X_train,y_train)
```

```
C:\Users\vedan\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.9_qbz5n2  
kfra8p0\LocalCache\local-packages\Python39\site-  
packages\sklearn\utils\validation.py:1300: DataConversionWarning: A column-  
vector y was passed when a 1d array was expected. Please change the shape of y  
to (n_samples, ), for example using ravel().  
y = column_or_1d(y, warn=True)
```

```
[11]: GaussianNB()
```

```
[12]: y_pred = gnb.predict(X_test)  
y_pred
```

```
[12]: array(['Iris-versicolor', 'Iris-setosa', 'Iris-virginica',  
            'Iris-versicolor', 'Iris-versicolor', 'Iris-setosa',  
            'Iris-versicolor', 'Iris-virginica', 'Iris-versicolor',  
            'Iris-versicolor', 'Iris-virginica', 'Iris-setosa', 'Iris-setosa',  
            'Iris-setosa', 'Iris-setosa', 'Iris-versicolor', 'Iris-virginica',  
            'Iris-versicolor', 'Iris-versicolor', 'Iris-virginica',  
            'Iris-setosa', 'Iris-virginica', 'Iris-setosa', 'Iris-virginica',  
            'Iris-virginica', 'Iris-virginica', 'Iris-virginica',  
            'Iris-virginica', 'Iris-setosa', 'Iris-setosa', 'Iris-setosa',  
            'Iris-setosa', 'Iris-versicolor', 'Iris-setosa', 'Iris-setosa',  
            'Iris-virginica', 'Iris-versicolor', 'Iris-setosa'], dtype='<U15')
```

```
[13]: from sklearn.metrics import accuracy_score , confusion_matrix , precision_score  
↳ , recall_score , classification_report
```

```
[14]: accuracy = accuracy_score(y_test,y_pred)  
accuracy
```

```
[14]: 1.0
```


```
[15]: cm = confusion_matrix(y_test,y_pred)  
cm
```

```
[15]: array([[15,  0,  0],  
          [ 0, 11,  0],  
          [ 0,  0, 12]], dtype=int64)
```



```
[16]: import seaborn as sns
import matplotlib.pyplot as plt
```

```
[17]: sns.heatmap(cm , annot=True , cbar=False , cmap='Oranges' ,
                xticklabels = ['Iris-setosa', 'Iris-versicolor', 'Iris-virginica'] ,
                yticklabels=['Iris-setosa', 'Iris-versicolor', 'Iris-virginica'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```



```
[18]: precision_score(y_test,y_pred , average='macro')
```

```
[18]: 1.0
```

```
[19]: recall_score(y_test,y_pred,average='macro')
```

```
[19]: 1.0
```

```
[20]: error_rate = 1 - accuracy
      error_rate
```

```
[20]: 0.0
```

```
[ ]:
```

final-ds-8

May 2, 2024

1 8) Data Visualization I

1. Use the inbuilt dataset 'titanic'. The dataset contains 891 rows and contains information about the passengers who boarded the unfortunate Titanic ship. Use the Seaborn library to see if we can find any patterns in the data.
2. Write a code to check how the price of the ticket (column name: 'fare') for each passenger is distributed by plotting a histogram

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: df = pd.read_csv('Titanic.csv')
df
```

```
[2]: PassengerId  Survived  Pclass  \
0               1         0       3
1               2         1       1
2               3         1       3
3               4         1       1
4               5         0       3
..            ...         ...     ...
886            887         0       2
887            888         1       1
888            889         0       3
889            890         1       1
890            891         0       3
```

```

                                Name    Sex  Age  SibSp  \
0          Braund, Mr. Owen Harris   male  22.0      1
1  Cumings, Mrs. John Bradley (Florence Briggs Th... female  38.0      1
2          Heikkinen, Miss. Laina   female  26.0      0
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)   female  35.0      1
4          Allen, Mr. William Henry   male  35.0      0
..            ...         ...     ...
886          Montvila, Rev. Juozas   male  27.0      0
887          Graham, Miss. Margaret Edith   female  19.0      0
888  Johnston, Miss. Catherine Helen "Carrie"   female   NaN      1
```

| | | | | | |
|-----|--|-----------------------|------|------|---|
| 889 | | Behr, Mr. Karl Howell | male | 26.0 | 0 |
| 890 | | Dooley, Mr. Patrick | male | 32.0 | 0 |

| | Parch | Ticket | Fare | Cabin | Embarked |
|-----|-------|------------------|---------|-------|----------|
| 0 | 0 | A/5 21171 | 7.2500 | NaN | S |
| 1 | 0 | PC 17599 | 71.2833 | C85 | C |
| 2 | 0 | STON/O2. 3101282 | 7.9250 | NaN | S |
| 3 | 0 | 113803 | 53.1000 | C123 | S |
| 4 | 0 | 373450 | 8.0500 | NaN | S |
| .. | ... | ... | ... | ... | ... |
| 886 | 0 | 211536 | 13.0000 | NaN | S |
| 887 | 0 | 112053 | 30.0000 | B42 | S |
| 888 | 2 | W./C. 6607 | 23.4500 | NaN | S |
| 889 | 0 | 111369 | 30.0000 | C148 | C |
| 890 | 0 | 370376 | 7.7500 | NaN | Q |

[891 rows x 12 columns]

```
[3]: df.isnull().sum()
```

```
[3]: PassengerId      0
Survived             0
Pclass              0
Name                0
Sex                 0
Age                177
SibSp              0
Parch              0
Ticket             0
Fare               0
Cabin              687
Embarked           2
dtype: int64
```

```
[4]: df.shape
```

```
[4]: (891, 12)
```

```
[8]: df.drop(columns='Cabin' , inplace=True)
```

```
[9]: df
```

```
[9]:   PassengerId  Survived  Pclass  \
0             1         0       3
1             2         1       1
2             3         1       3
3             4         1       1
```

```

4          5          0          3
..      ...      ...      ...
886      887          0          2
887      888          1          1
888      889          0          3
889      890          1          1
890      891          0          3

```

```

                                Name      Sex  Age  SibSp  \
0                Braund, Mr. Owen Harris   male  22.0    1
1  Cumings, Mrs. John Bradley (Florence Briggs Th... female  38.0    1
2                Heikkinen, Miss. Laina   female  26.0    0
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)   female  35.0    1
4                Allen, Mr. William Henry   male  35.0    0
..
886                Montvila, Rev. Juozas   male  27.0    0
887                Graham, Miss. Margaret Edith   female  19.0    0
888      Johnston, Miss. Catherine Helen "Carrie"   female   NaN    1
889                Behr, Mr. Karl Howell   male  26.0    0
890                Dooley, Mr. Patrick   male  32.0    0

```

```

      Parch      Ticket    Fare Embarked
0         0      A/5 21171    7.2500        S
1         0      PC 17599   71.2833        C
2         0  STON/O2. 3101282    7.9250        S
3         0      113803   53.1000        S
4         0      373450    8.0500        S
..      ...
886        0      211536   13.0000        S
887        0      112053   30.0000        S
888        2      W./C. 6607   23.4500        S
889        0      111369   30.0000        C
890        0      370376    7.7500        Q

```

[891 rows x 11 columns]

```
[10]: df.dropna(inplace=True)
```

```
[11]: df.shape
```

```
[11]: (712, 11)
```

```
[12]: df
```

```

[12]:      PassengerId  Survived  Pclass  \
0             1         0         3
1             2         1         1

```

| | | | |
|-----|-----|-----|-----|
| 2 | 3 | 1 | 3 |
| 3 | 4 | 1 | 1 |
| 4 | 5 | 0 | 3 |
| .. | ... | ... | ... |
| 885 | 886 | 0 | 3 |
| 886 | 887 | 0 | 2 |
| 887 | 888 | 1 | 1 |
| 889 | 890 | 1 | 1 |
| 890 | 891 | 0 | 3 |

| | Name | Sex | Age | SibSp | \ |
|-----|--|--------|------|-------|---|
| 0 | Braund, Mr. Owen Harris | male | 22.0 | 1 | |
| 1 | Cummings, Mrs. John Bradley (Florence Briggs Th... | female | 38.0 | 1 | |
| 2 | Heikkinen, Miss. Laina | female | 26.0 | 0 | |
| 3 | Futrelle, Mrs. Jacques Heath (Lily May Peel) | female | 35.0 | 1 | |
| 4 | Allen, Mr. William Henry | male | 35.0 | 0 | |
| .. | ... | ... | ... | ... | |
| 885 | Rice, Mrs. William (Margaret Norton) | female | 39.0 | 0 | |
| 886 | Montvila, Rev. Juozas | male | 27.0 | 0 | |
| 887 | Graham, Miss. Margaret Edith | female | 19.0 | 0 | |
| 889 | Behr, Mr. Karl Howell | male | 26.0 | 0 | |
| 890 | Dooley, Mr. Patrick | male | 32.0 | 0 | |

| | Parch | Ticket | Fare | Embarked |
|-----|-------|------------------|---------|----------|
| 0 | 0 | A/5 21171 | 7.2500 | S |
| 1 | 0 | PC 17599 | 71.2833 | C |
| 2 | 0 | STON/O2. 3101282 | 7.9250 | S |
| 3 | 0 | 113803 | 53.1000 | S |
| 4 | 0 | 373450 | 8.0500 | S |
| .. | ... | ... | ... | ... |
| 885 | 5 | 382652 | 29.1250 | Q |
| 886 | 0 | 211536 | 13.0000 | S |
| 887 | 0 | 112053 | 30.0000 | S |
| 889 | 0 | 111369 | 30.0000 | C |
| 890 | 0 | 370376 | 7.7500 | Q |

[712 rows x 11 columns]

```
[13]: df.columns
```

```
[13]: Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',
          'Parch', 'Ticket', 'Fare', 'Embarked'],
          dtype='object')
```

```
[14]: df.drop(columns='PassengerId' , inplace = True)
```

```
[15]: df
```

```
[15]:
```


| | Survived | Pclass | Name \ |
|-----|----------|--------|---|
| 0 | 0 | 3 | Braund, Mr. Owen Harris |
| 1 | 1 | 1 | Cumings, Mrs. John Bradley (Florence Briggs Th... |
| 2 | 1 | 3 | Heikkinen, Miss. Laina |
| 3 | 1 | 1 | Futrelle, Mrs. Jacques Heath (Lily May Peel) |
| 4 | 0 | 3 | Allen, Mr. William Henry |
| .. | ... | ... | ... |
| 885 | 0 | 3 | Rice, Mrs. William (Margaret Norton) |
| 886 | 0 | 2 | Montvila, Rev. Juozas |
| 887 | 1 | 1 | Graham, Miss. Margaret Edith |
| 889 | 1 | 1 | Behr, Mr. Karl Howell |
| 890 | 0 | 3 | Dooley, Mr. Patrick |

| | Sex | Age | SibSp | Parch | Ticket | Fare | Embarked |
|-----|--------|------|-------|-------|------------------|---------|----------|
| 0 | male | 22.0 | 1 | 0 | A/5 21171 | 7.2500 | S |
| 1 | female | 38.0 | 1 | 0 | PC 17599 | 71.2833 | C |
| 2 | female | 26.0 | 0 | 0 | STON/O2. 3101282 | 7.9250 | S |
| 3 | female | 35.0 | 1 | 0 | 113803 | 53.1000 | S |
| 4 | male | 35.0 | 0 | 0 | 373450 | 8.0500 | S |
| .. | ... | ... | ... | ... | ... | ... | ... |
| 885 | female | 39.0 | 0 | 5 | 382652 | 29.1250 | Q |
| 886 | male | 27.0 | 0 | 0 | 211536 | 13.0000 | S |
| 887 | female | 19.0 | 0 | 0 | 112053 | 30.0000 | S |
| 889 | male | 26.0 | 0 | 0 | 111369 | 30.0000 | C |
| 890 | male | 32.0 | 0 | 0 | 370376 | 7.7500 | Q |

[712 rows x 10 columns]


2 Survival by Gender

```
[17]: # Bar plot showing the count of survivors by gender
sns.countplot(x="Sex", hue="Survived", data=df)
plt.title("Survival by Gender")
plt.show()
```



3 Survival by Passenger Class


```
[18]: # Bar plot showing the count of survivors by passenger class
sns.countplot(x="Pclass", hue="Survived", data=df)
plt.title("Survival by Passenger Class")
plt.show()
```





4 Age Distribution

```
[19]: # Distribution plot showing age distribution
sns.histplot(df["Age"], kde=True)
plt.title("Age Distribution")
plt.show()


# Violin plot showing age distribution by survival
sns.violinplot(x="Survived", y="Age", data=df)
plt.title("Age Distribution by Survival")
plt.show()
```






5 Fare Distribution

```
[24]: sns.histplot(df['Fare'], kde=True) # kde=True plots a Kernel Density Estimate,
      ↪ (smoothed curve)
      plt.xlabel('Fare')
      plt.ylabel('Count')
      plt.title('Distribution of Ticket Prices (Fare) for Titanic Passengers')
      plt.show()
```



6 Survival by Embark Point

```
[28]: # Bar plot showing survival by embarkation point
sns.countplot(x="Embarked", hue="Survived", data=df)
plt.title("Survival by Embarkation Point")
plt.show()
#C = Cherbourg, Q = Queenstown, S = Southampton
```



[]:

final-ds-9

May 2, 2024

1 9 Data Visualization II

1. Use the inbuilt dataset 'titanic' as used in the above problem. Plot a box plot for distribution of age with respect to each gender along with the information about whether they survived or not. (Column names : 'sex' and 'age')
2. Write observations on the inference from the above statistics.

```
[1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[3]: df = pd.read_csv('Titanic.csv')
df
```

```
[3]:
```

| | PassengerId | Survived | Pclass | \ |
|-----|-------------|----------|--------|---|
| 0 | 1 | 0 | 3 | |
| 1 | 2 | 1 | 1 | |
| 2 | 3 | 1 | 3 | |
| 3 | 4 | 1 | 1 | |
| 4 | 5 | 0 | 3 | |
| .. | ... | ... | ... | |
| 886 | 887 | 0 | 2 | |
| 887 | 888 | 1 | 1 | |
| 888 | 889 | 0 | 3 | |
| 889 | 890 | 1 | 1 | |
| 890 | 891 | 0 | 3 | |

| | Name | Sex | Age | SibSp | \ |
|-----|---|--------|------|-------|---|
| 0 | Braund, Mr. Owen Harris | male | 22.0 | 1 | |
| 1 | Cumings, Mrs. John Bradley (Florence Briggs Th... | female | 38.0 | 1 | |
| 2 | Heikkinen, Miss. Laina | female | 26.0 | 0 | |
| 3 | Futrelle, Mrs. Jacques Heath (Lily May Peel) | female | 35.0 | 1 | |
| 4 | Allen, Mr. William Henry | male | 35.0 | 0 | |
| .. | ... | ... | ... | ... | |
| 886 | Montvila, Rev. Juozas | male | 27.0 | 0 | |
| 887 | Graham, Miss. Margaret Edith | female | 19.0 | 0 | |
| 888 | Johnston, Miss. Catherine Helen "Carrie" | female | NaN | 1 | |
| 889 | Behr, Mr. Karl Howell | male | 26.0 | 0 | |


890 Dooley, Mr. Patrick male 32.0 0

| | Parch | Ticket | Fare | Cabin | Embarked |
|-----|-------|------------------|---------|-------|----------|
| 0 | 0 | A/5 21171 | 7.2500 | NaN | S |
| 1 | 0 | PC 17599 | 71.2833 | C85 | C |
| 2 | 0 | STON/O2. 3101282 | 7.9250 | NaN | S |
| 3 | 0 | 113803 | 53.1000 | C123 | S |
| 4 | 0 | 373450 | 8.0500 | NaN | S |
| .. | .. | ... | ... | ... | ... |
| 886 | 0 | 211536 | 13.0000 | NaN | S |
| 887 | 0 | 112053 | 30.0000 | B42 | S |
| 888 | 2 | W./C. 6607 | 23.4500 | NaN | S |
| 889 | 0 | 111369 | 30.0000 | C148 | C |
| 890 | 0 | 370376 | 7.7500 | NaN | Q |

[891 rows x 12 columns]

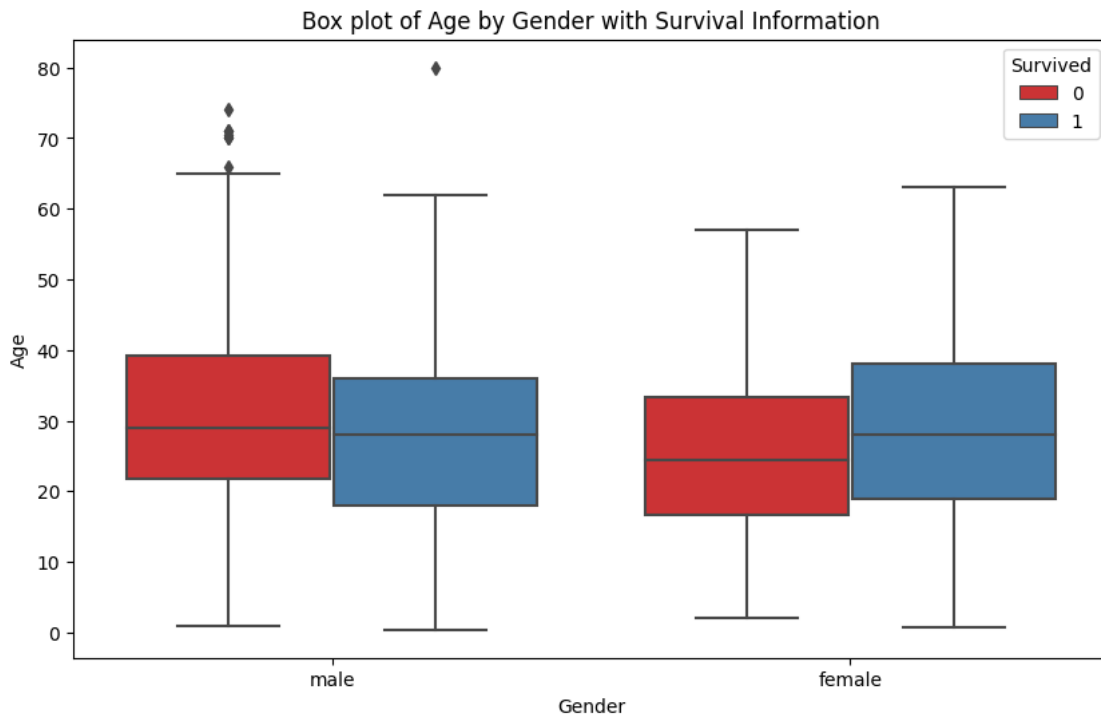
```
[6]: sns.boxplot(x='Sex' , y='Age' , data=df , hue = 'Survived')
```

```
[6]: <Axes: xlabel='Sex', ylabel='Age'>
```



```
[10]: plt.figure(figsize=(10, 6))
sns.boxplot(data=df, x='Sex', y='Age', hue='Survived', palette='Set1')

plt.title('Box plot of Age by Gender with Survival Information')
plt.xlabel('Gender')
plt.ylabel('Age')
plt.legend(title='Survived', loc='upper right')
plt.show()
```



2 Matplotlib

Matplotlib is a comprehensive plotting library for creating static, interactive, and animated visualizations in Python.

3 Seaborn

Seaborn is a data visualization library built on top of Matplotlib, designed to simplify complex statistical plots and make them more visually appealing

4 Boxplot

A boxplot (also known as a box-and-whisker plot) is a statistical visualization that displays the distribution of a dataset through its quartiles. It shows the median, upper and lower quartiles, and

potential outliers, using a box to represent the interquartile range and “whiskers” to indicate the range of the data

[]:

final-ds-10

May 2, 2024

1 Data Visualization III

Download the Iris flower dataset or any other dataset into a DataFrame. (e.g., <https://archive.ics.uci.edu/ml/datasets/Iris>). Scan the dataset and give the inference as: 1. List down the features and their types (e.g., numeric, nominal) available in the dataset. 2. Create a histogram for each feature in the dataset to illustrate the feature distributions. 3. Create a boxplot for each feature in the dataset. 4. Compare distributions and identify outliers.

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: df = pd.read_csv('Iris (1).csv')
df
```

```
[2]:      Id  SepalLengthCm  SepalWidthCm  PetalLengthCm  PetalWidthCm  \
0      1             5.1             3.5             1.4             0.2
1      2             4.9             3.0             1.4             0.2
2      3             4.7             3.2             1.3             0.2
3      4             4.6             3.1             1.5             0.2
4      5             5.0             3.6             1.4             0.2
..  ...             ...             ...             ...             ...
145 146             6.7             3.0             5.2             2.3
146 147             6.3             2.5             5.0             1.9
147 148             6.5             3.0             5.2             2.0
148 149             6.2             3.4             5.4             2.3
149 150             5.9             3.0             5.1             1.8
```

```
      Species
0      Iris-setosa
1      Iris-setosa
2      Iris-setosa
3      Iris-setosa
4      Iris-setosa
..  ...
145  Iris-virginica
146  Iris-virginica
147  Iris-virginica
```

```
148 Iris-virginica
149 Iris-virginica
```

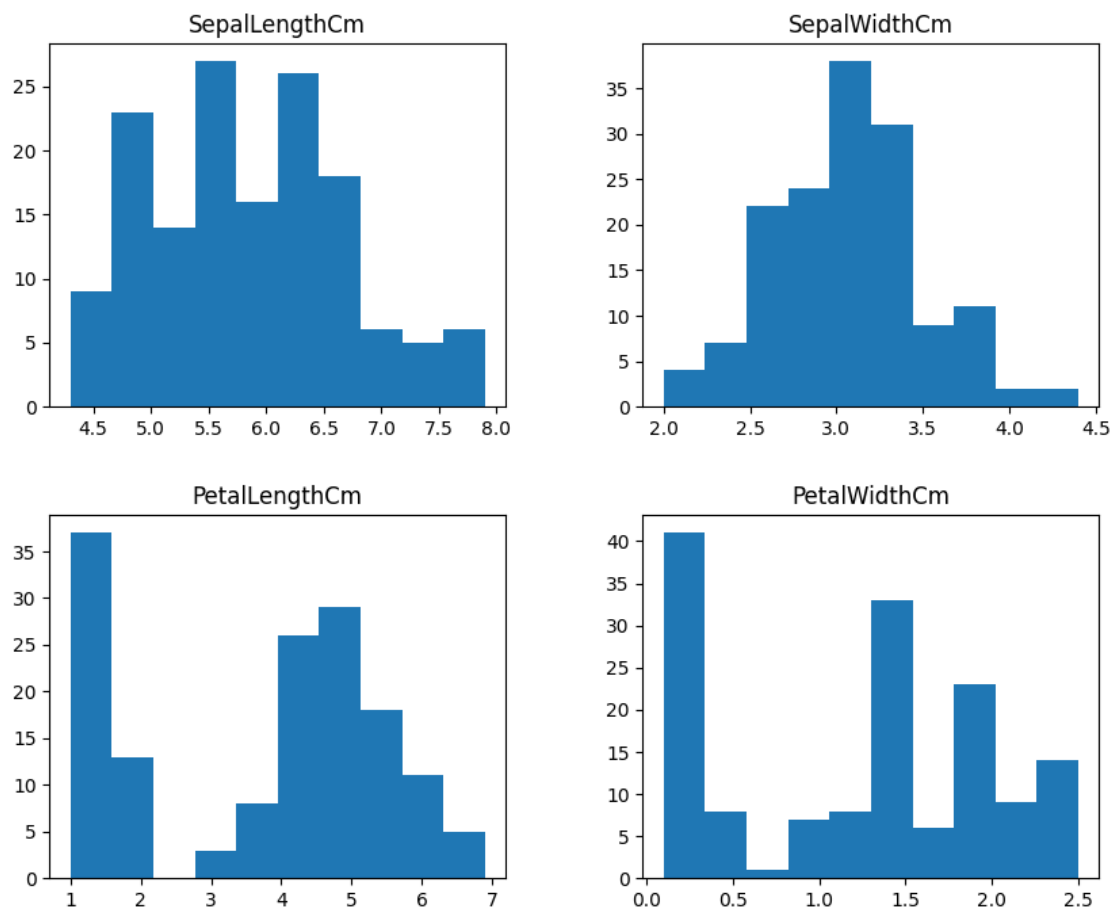
```
[150 rows x 6 columns]
```

```
[3]: df.drop(columns='Id' , inplace=True)
```

```
[4]: df.dtypes
```


```
[4]: SepalLengthCm    float64
     SepalWidthCm     float64
     PetalLengthCm    float64
     PetalWidthCm     float64
     Species         object
     dtype: object
```

```
[5]: # Create a histogram for each numeric feature
     df.hist(bins=10, figsize=(10, 8), grid=False)
     plt.show()
```



```
[6]: fig, axes = plt.subplots(nrows=2,ncols=2,figsize = (9,9))
axes = axes.flatten()
ax = sns.histplot(x='SepalLengthCm', data=df, bins=10,ax=axes[0])

ax = sns.histplot(x='SepalWidthCm', data=df, bins=10,ax=axes[1])
ax = sns.histplot(x='PetalLengthCm', data=df, bins=10,ax=axes[2])
ax = sns.histplot(x='PetalWidthCm', data=df, bins=10,ax=axes[3])
```



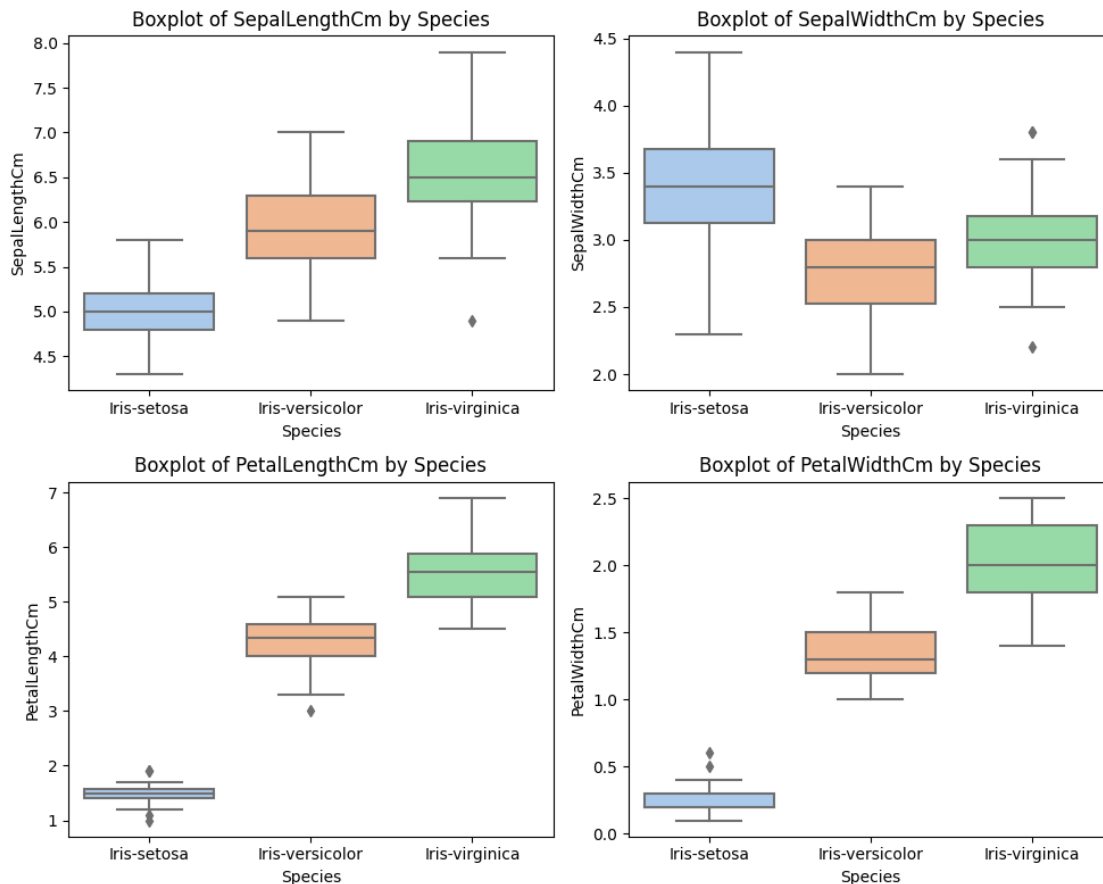
```
[7]: col = df.columns
col[:-1]
```

```
[7]: Index(['SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm'],
dtype='object')
```

```
[8]: plt.figure(figsize=(10, 8))
for i, feature in enumerate(col[:-1]):
    # Create a 2x2 grid of subplots
    plt.subplot(2, 2, i + 1) # 2 rows, 2 columns
    sns.boxplot(x='Species', y=feature, data=df, palette="pastel")

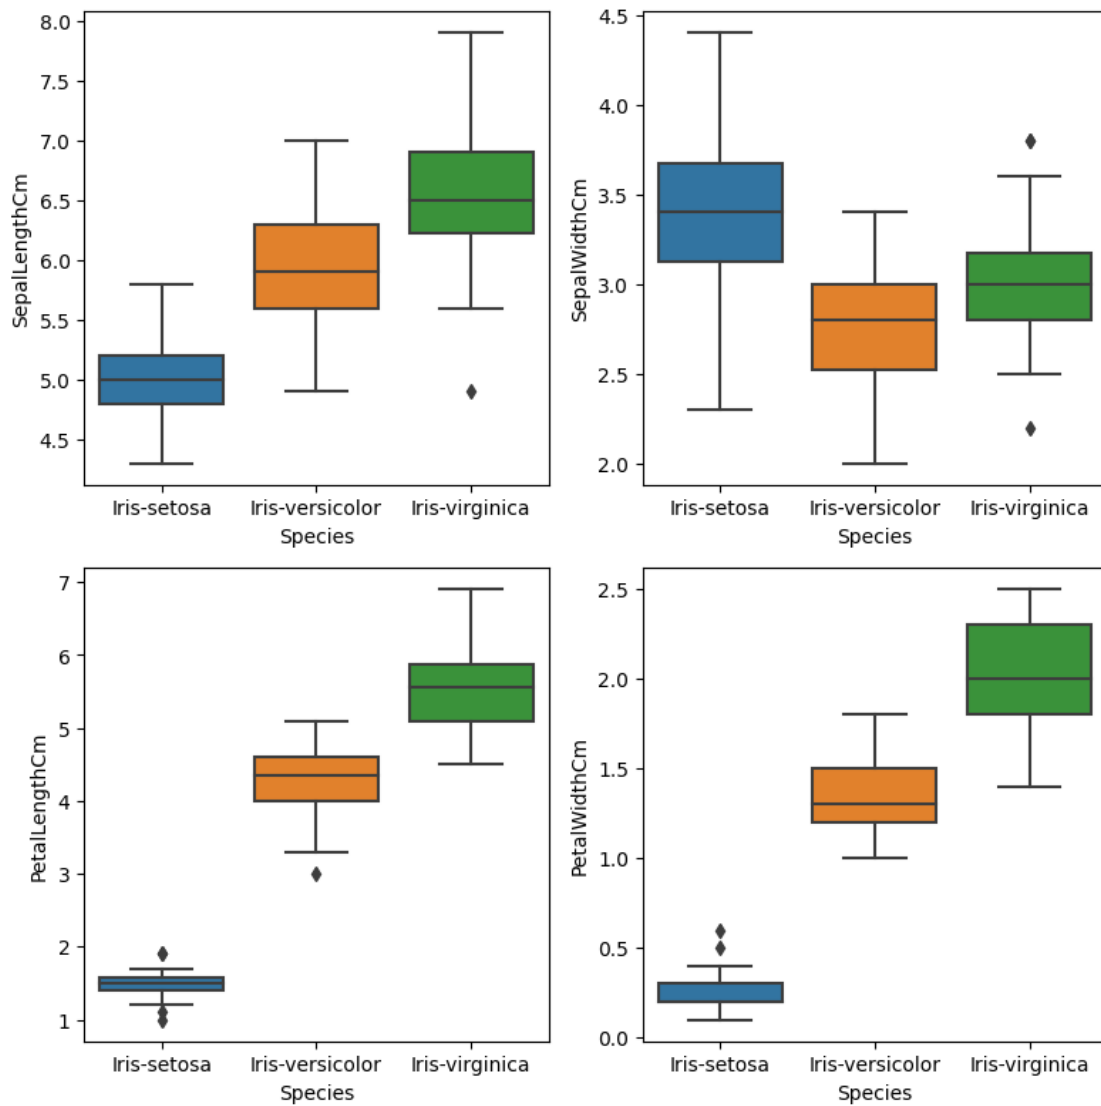
    # Title and rotation
    plt.title(f'Boxplot of {feature} by Species')
    # plt.xticks(rotation=45)

# Adjust layout to avoid overlap
plt.tight_layout() # This adjusts spacing between plots
plt.show()
```



```
[9]: fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(8, 8))
axes = axes.flatten()
```

```
sns.boxplot(x='Species' , y = 'SepalLengthCm' ,data=df, ax=axes[0])
sns.boxplot(x='Species' , y = 'SepalWidthCm' ,data=df, ax=axes[1])
sns.boxplot(x='Species' , y = 'PetalLengthCm' ,data=df, ax=axes[2])
sns.boxplot(x='Species' , y = 'PetalWidthCm' ,data=df, ax=axes[3])
plt.tight_layout()
```



[]:

final-ds-7-1

May 2, 2024

```
[ ]: import nltk
nltk.download('punkt')
nltk.download('wordnet')
nltk.download('averaged_perceptron_tagger')
nltk.download('stopwords')
from nltk import sent_tokenize
from nltk import word_tokenize
from nltk.corpus import stopwords
```

```
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data]   Package punkt is already up-to-date!
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data]   Package wordnet is already up-to-date!
[nltk_data] Downloading package averaged_perceptron_tagger to
[nltk_data]   /root/nltk_data...
[nltk_data]   Package averaged_perceptron_tagger is already up-to-
[nltk_data]   date!
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Package stopwords is already up-to-date!
```

```
[ ]: text = 'Researchers observed that the children were playing with their toys in_
↳the park while their parents were discussing various topics nearby. As the_
↳afternoon grew warm, families began preparing for a picnic, spreading_
↳blankets and setting up food. The wind picked up, rustling the leaves, and_
↳everyone enjoyed the calm and peaceful atmosphere.'
```

```
[ ]: tokens_sents = nltk.sent_tokenize(text)
print(tokens_sents)
```

```
['Researchers observed that the children were playing with their toys in the
park while their parents were discussing various topics nearby.', 'As the
afternoon grew warm, families began preparing for a picnic, spreading blankets
and setting up food.', 'The wind picked up, rustling the leaves, and everyone
enjoyed the calm and peaceful atmosphere.']
```

```
[ ]: tokens_words = nltk.word_tokenize(text)
print(tokens_words)
```

```
['Researchers', 'observed', 'that', 'the', 'children', 'were', 'playing',
'with', 'their', 'toys', 'in', 'the', 'park', 'while', 'their', 'parents',
'were', 'discussing', 'various', 'topics', 'nearby', '.', 'As', 'the',
'afternoon', 'grew', 'warm', ',', 'families', 'began', 'preparing', 'for', 'a',
'picnic', ',', 'spreading', 'blankets', 'and', 'setting', 'up', 'food', '.',
'The', 'wind', 'picked', 'up', ',', 'rustling', 'the', 'leaves', ',', 'and',
'everyone', 'enjoyed', 'the', 'calm', 'and', 'peaceful', 'atmosphere', '.']
```

```
[ ]: sw_nltk = stopwords.words('english')
      print((sw_nltk))
```

```
['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're",
"you've", "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he',
'him', 'his', 'himself', 'she', "she's", 'her', 'hers', 'herself', 'it', "it's",
'its', 'itself', 'they', 'them', 'their', 'theirs', 'themselves', 'what',
'which', 'who', 'whom', 'this', 'that', "that'll", 'these', 'those', 'am', 'is',
'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had', 'having',
'do', 'does', 'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or',
'because', 'as', 'until', 'while', 'of', 'at', 'by', 'for', 'with', 'about',
'against', 'between', 'into', 'through', 'during', 'before', 'after', 'above',
'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'over', 'under',
'again', 'further', 'then', 'once', 'here', 'there', 'when', 'where', 'why',
'how', 'all', 'any', 'both', 'each', 'few', 'more', 'most', 'other', 'some',
'such', 'no', 'nor', 'not', 'only', 'own', 'same', 'so', 'than', 'too', 'very',
's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'now',
'd', 'll', 'm', 'o', 're', 've', 'y', 'ain', 'aren', "aren't", 'couldn',
"couldn't", 'didn', "didn't", 'doesn', "doesn't", 'hadn', "hadn't", 'hasn',
"hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mightn', "mightn't",
'mustn', "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn',
"shouldn't", 'wasn', "wasn't", 'weren', "weren't", 'won', "won't", 'wouldn',
"wouldn't"]
```

```
[ ]: # from nltk.corpus import stopwords
      # stop_words = set(stopwords.words('english'))
      # stop_words
```

```
[ ]: words = [i for i in text.split() if i.lower() not in sw_nltk]
      # new_text = " ".join(words)
      # print(new_text)
      words
```

```
[ ]: ['Researchers',
      'observed',
      'children',
      'playing',
      'toys',
      'park',
```



```

'parents',
'discussing',
'various',
'topics',
'nearby.',
'afternoon',
'grew',
'warm,',
'families',
'began',
'preparing',
'picnic,',
'spreading',
'blankets',
'setting',
'food.',
'wind',
'picked',
'up,',
'rustling',
'leaves,',
'everyone',
'enjoyed',
'calm',
'peaceful',
'atmosphere.']

```

```
[ ]: from nltk.stem import PorterStemmer
```

```
[ ]: stem=[]
for i in words:
    ps = PorterStemmer()
    stem_word= ps.stem(i)
    stem.append(stem_word)
print(stem)
```

```

['research', 'observ', 'children', 'play', 'toy', 'park', 'parent', 'discuss',
'variou', 'topic', 'nearby.', 'afternoon', 'grew', 'warm,', 'famili', 'began',
'prepar', 'picnic,', 'spread', 'blanket', 'set', 'food.', 'wind', 'pick', 'up,',
'rustl', 'leaves,', 'everyon', 'enjoy', 'calm', 'peac', 'atmosphere.']

```

```
[ ]: import nltk
from nltk.stem import WordNetLemmatizer
lemmatizer = WordNetLemmatizer()
```

```
[ ]: leme=[]
for i in words:
```

```

    lemetized_word=lemmatizer.lemmatize(i)
    leme.append(lemetized_word)
print(leme)

new_text = " ".join(leme)
print(new_text)

```

```

['Researchers', 'observed', 'child', 'playing', 'toy', 'park', 'parent',
'discussing', 'various', 'topic', 'nearby.', 'afternoon', 'grew', 'warm,',
'family', 'began', 'preparing', 'picnic,', 'spreading', 'blanket', 'setting',
'food.', 'wind', 'picked', 'up,', 'rustling', 'leaves,', 'everyone', 'enjoyed',
'calm', 'peaceful', 'atmosphere.']

```

Researchers observed child playing toy park parent discussing various topic nearby. afternoon grew warm, family began preparing picnic, spreading blanket setting food. wind picked up, rustling leaves, everyone enjoyed calm peaceful atmosphere.

```
[ ]: print("Parts of Speech: ",nltk.pos_tag(leme))
```

```

Parts of Speech: [('Researchers', 'NNS'), ('observed', 'VBD'), ('child', 'JJ'),
('playing', 'VBG'), ('toy', 'NN'), ('park', 'NN'), ('parent', 'NN'),
('discussing', 'VBG'), ('various', 'JJ'), ('topic', 'NN'), ('nearby.', 'NN'),
('afternoon', 'NN'), ('grew', 'VBD'), ('warm,', 'JJ'), ('family', 'NN'),
('began', 'VBD'), ('preparing', 'VBG'), ('picnic,', 'NN'), ('spreading', 'VBG'),
('blanket', 'NN'), ('setting', 'VBG'), ('food.', 'JJ'), ('wind', 'NN'),
('picked', 'VBD'), ('up,', 'JJ'), ('rustling', 'VBG'), ('leaves,', 'NN'),
('everyone', 'NN'), ('enjoyed', 'VBD'), ('calm', 'JJ'), ('peaceful', 'JJ'),
('atmosphere.', 'NN')]

```

```
[ ]: from sklearn.feature_extraction.text import TfidfVectorizer
```

```
[ ]: vectorizer = TfidfVectorizer()
      tfidf_matrix = vectorizer.fit_transform([new_text])
      # Get the feature names (terms)
      feature_names = vectorizer.get_feature_names_out()
      # Print the document-term matrix
      print("Document 1")
      for j, feature in enumerate(feature_names):
          if tfidf_matrix[0,j] > 0:
              print(" ", feature, ":", tfidf_matrix[0,j])

```

Document 1

```

afternoon : 0.17677669529663687
atmosphere : 0.17677669529663687
began : 0.17677669529663687
blanket : 0.17677669529663687
calm : 0.17677669529663687
child : 0.17677669529663687

```

discussing : 0.17677669529663687
enjoyed : 0.17677669529663687
everyone : 0.17677669529663687
family : 0.17677669529663687
food : 0.17677669529663687
grew : 0.17677669529663687
leaves : 0.17677669529663687
nearby : 0.17677669529663687
observed : 0.17677669529663687
parent : 0.17677669529663687
park : 0.17677669529663687
peaceful : 0.17677669529663687
picked : 0.17677669529663687
picnic : 0.17677669529663687
playing : 0.17677669529663687
preparing : 0.17677669529663687
researchers : 0.17677669529663687
rustling : 0.17677669529663687
setting : 0.17677669529663687
spreading : 0.17677669529663687
topic : 0.17677669529663687
toy : 0.17677669529663687
up : 0.17677669529663687
various : 0.17677669529663687
warm : 0.17677669529663687
wind : 0.17677669529663687

[]:

untitled

May 2, 2024

1 NLTK

The Natural Language Toolkit (nltk) is a comprehensive library in Python designed for natural language processing (NLP) tasks.

It provides tools for text processing, such as tokenization, parsing, stemming, and machine learning algorithms for text analysis

2 punkt

used for tokenizing into words and sentences

3 Stopwords

Stopwords are common words in a language (such as “and,” “the,” “is,” etc.) that are typically ignored or removed in text processing and analysis because they carry less significant meaning. These words are usually filtered out to focus on the more meaningful terms that contribute to understanding the context or content of the text.

4 POS Tagging

assigning labels to words in a sentence to identify their grammatical roles, such as nouns, verbs, adjectives, or adverbs.

JJ -> ADJECTIVE

NN - > NOUN

VB - > VERB

5 Stemming

Stemming is the process of reducing words to their root form by removing suffixes or prefixes, like turning “running” into “run.” It’s used to standardize text for easier analysis and search

After stemming its not necessary that the word will have its meaning in english dictionary

Its used , when we dont want to show output to user and its complexity is less

6 Lemmatization

Lemmatization is the process of converting words to their base or dictionary form, known as the lemma, by considering the context and grammatical rules

After Lemmatization word have meaning in english dictionary

Its used when we want to provide answer / output to user , and its complexity is high

7 TFID Vectorizer

Full Form : Term Frequency-Inverse Document Frequency

It combines two metrics:

Term Frequency (TF): Measures how often a word appears in a document. A higher frequency indicates greater relevance within that document.

Inverse Document Frequency (IDF): Measures how common or rare a word is across a set of documents. A higher IDF indicates that the word is rare, suggesting it carries more distinctive information.

8 $TF = \text{frequency of term 't' in document 'd'} / \text{total terms in 'd'}$

9 $IDF = \log_{10} (\text{total number of documents} / \text{total documents with term 't'})$

[]: